

## Hyperparameter Tuning for kNN for Predicting Heart Disease

```
from google.colab import files
upload = files.upload()
```

No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.  
Saving heart (2).csv to heart (2) (1).csv

```
import pandas as pd
import seaborn as sns
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import GridSearchCV #cross validation to select the optimal params for knn
# k distance will be considered under a given set of values.
```

```
df = pd.read_csv('heart (2).csv')
```

```
print(df.info())
print(df.shape)
print(df.head())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	age	303 non-null	int64
1	sex	303 non-null	int64
2	cp	303 non-null	int64
3	trestbps	303 non-null	int64
4	chol	303 non-null	int64
5	fbs	303 non-null	int64
6	restecg	303 non-null	int64
7	thalach	303 non-null	int64
8	exang	303 non-null	int64
9	oldpeak	303 non-null	float64
10	slope	303 non-null	int64
11	ca	303 non-null	int64
12	thal	303 non-null	int64
13	target	303 non-null	int64

```
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
None
```

```
(303, 14)
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
df.isnull().sum()
```


```
age      0
sex      0
cp       0
trestbps 0
chol     0
fbs      0
restecg  0
thalach  0
exang    0
oldpeak  0
slope    0
```

```
ca      0
thal    0
target  0
dtype: int64
```

```
sns.countplot(df['target'])
```

```
/usr/local/lib/python3.8/dist-packages/seaborn/_decorators.py:36: FutureWarning: Pass the following variables as keyword arguments: {'x': 'target', 'y': 'count'}. This warning will be removed in a future version of seaborn.
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f029334b670>
```



```
x=df.drop(columns=['target'])
y=df['target']
knn=KNeighborsClassifier()
train_X,test_X,train_Y,test_Y=train_test_split(x,y,train_size=0.2, random_state=100)
```

train a new knn model using gscv

```
knn.fit(train_X,train_Y)
KNeighborsClassifier()
```

```
y_pred=knn.predict(test_X)
```

```
print(classification_report(test_Y,y_pred))
```

	precision	recall	f1-score	support
0	0.53	0.51	0.52	108
1	0.62	0.64	0.63	135
accuracy			0.58	243
macro avg	0.57	0.57	0.57	243
weighted avg	0.58	0.58	0.58	243

```
roc_auc_score(test_Y,y_pred)
```

```
0.5731481481481482
```

```
#set params using gscv/ set leaf size to 1 neigh to 10 and dist metric to 1,2 when 1 manhattan dis 2 euclidian dis. gscv to search the optima
leaf_size=list(range(1,15))
n_neighbors=list(range(1,10))
p=[1,2]
hyperparameters=dict(leaf_size=leaf_size,n_neighbors=n_neighbors,p=p)
```

```
knn_2 = KNeighborsClassifier()
clf = GridSearchCV(knn_2, hyperparameters, cv=10, scoring = 'roc_auc')
best_model = clf.fit(x,y)
```

```
print(hyperparameters)
```

```
{'leaf_size': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14], 'n_neighbors': [1, 2, 3, 4, 5, 6, 7, 8, 9], 'p': [1, 2]}
```

```
print(best_model.best_estimator_.get_params())
```

```
{'algorithm': 'auto', 'leaf_size': 9, 'metric': 'minkowski', 'metric_params': None, 'n_jobs': None, 'n_neighbors': 7, 'p': 1, 'weights':
```

```
print('Best leaf_size:', best_model.best_estimator_.get_params()['leaf_size'])
print('Best p:', best_model.best_estimator_.get_params()['p'])
print('Best n_neighbors:', best_model.best_estimator_.get_params()['n_neighbors'])
print('Best Score:', best_model.best_score_)
```

```
Best leaf_size: 9
Best p: 1
Best n_neighbors: 7
Best Score: 0.7483536683904332
```

```
x=df.drop(columns=['target'])
y=df['target']
knn=KNeighborsClassifier()
train_X,test_X,train_Y,test_Y=train_test_split(x,y,train_size=0.2, random_state=4)
```

```
knn.fit(train_X,train_Y)
```

```
KNeighborsClassifier()
```

```
y_pred=knn.predict(test_X)
```

```
print(classification_report(test_Y,y_pred))
```

	precision	recall	f1-score	support
0	0.53	0.66	0.59	106
1	0.67	0.54	0.60	137
accuracy			0.59	243
macro avg	0.60	0.60	0.59	243
weighted avg	0.61	0.59	0.59	243

```
roc_auc_score(test_Y,y_pred)
```

```
0.600261671946013
```